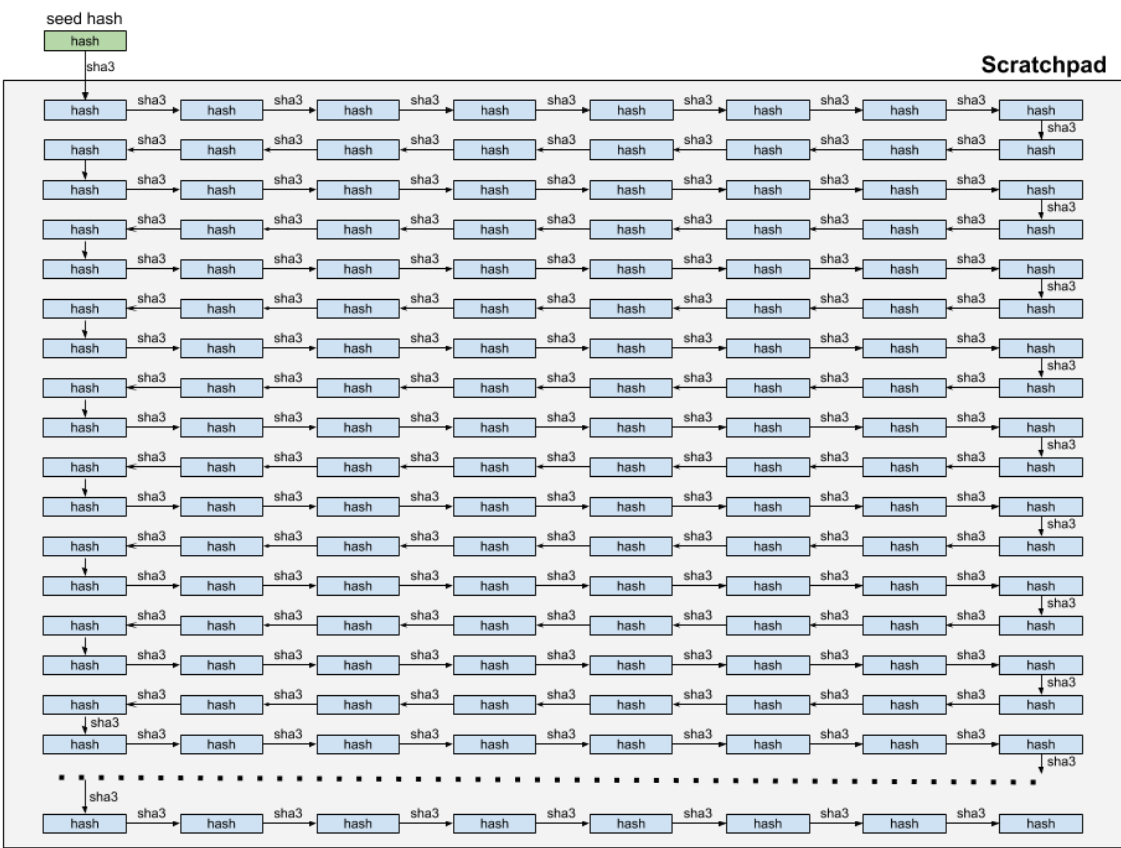


WildKeccak2 Scratchpad Generation Analysis



By calculating a scratchpad's items "on the fly", memory-hard algorithms may be optimized from resultant memory access operations. As such, the main advantages of a memory-hard PoW hash are eliminated. In this draft proposal, we present a theoretical analysis supporting the question of whether such operations would be reasonable for WildKeccak2 while using our current scratchpad generation scheme.

At this time of writing, our scratchpad calculation looks like the following:



Initially derived from seed, each next group of 32 bytes (shown as a hash for easier interpretation, let's call it "item") is generated from the previous group by calculating a sha3 hash. After all, we are getting a relatively big (starting from 500MB) solid memory block which is used as a scratchpad data for hashing in WildKeccak2. Each WildKeccak2 hash function call performs 7200 random read operations on this scratchpad.

With this approach, it is possible to have a light version of the scratchpad (the same idea as used in EthHash) which would be a trade-off between performance and memory consumption. For example, we can keep every 10-nth item and call them "bearing items". Then, during the hash calculation, every



Both graphics show the same sequence: the left side displayed in logarithmic scale, the right side displayed in linear. This made for an easier reading of difference from the very beginning of the graph (with a very small size of scratchpad, where all operations mostly fit in CPU's 1st level cache memory) compared to very last part, where the length of time appears to stop growing because nearly all memory access operations are addressed to global memory (DRAM).

Here, we can see in the red line that WK2 running on very fast 1st level cache memory takes about 1400ms and WK2 running on global memory is about 5000ms (all numbers measured against 10000 hash operations). So, the latency loss on accessing global memory is $5000 - 1400 = 3600$ ms for groupings of 7200 random read operations, with approximately 1/2 ms for each operation (still talking in terms of x10000 basis). And, the blue line, on x10000 basis, a SHA3 operation takes an average of 40 ms.

Following this logic, 1 random memory read cost (in a very rough approximation) is 1/80th of a SHA3 calculation coast. This leads us to believe that generating a WK2 scratchpad (that derives each subsequent item from the previous by only taking 1 hash) is a sensible solution because the cost to calculate at least one item is roughly x80 more then if read from global memory(normally need to calculate from 1 until 9 items).

